

# PyCUDA and GPUArray

See <https://documen.tician.de/pycuda/> for more documentation.

## Initialization

A few modules have to be loaded to initialize communication to the GPU:

```
import pycuda.driver as cuda
import pycuda.autoinit
```

The `pycuda.driver` module has methods that:

1. Allocate memory on the GPU (`cuda.mem_alloc()`)
2. Copy numpy arrays to the GPU (`cuda.memcpy_htod()`)
3. Execute kernels on the GPU described by CUDA code (see `compiler.SourceModule`)
4. Copy data from the GPU back to numpy arrays (`cuda.memcpy_dtoh()`).

## GPUArrays

The `pycuda.gpuarray.GPUArray` is particularly useful since they can associate a numpy array with an array on the device, handle transfers, and express array operations on the GPU in with numpy array syntax. E.g.:

```
import pycuda.gpuarray as gpuarray
import pycuda.driver as cuda
import pycuda.autoinit
import numpy

a_gpu = gpuarray.to_gpu(numpy.random.randn(4,4).astype(numpy.float32))
a_doubled = (2*a_gpu).get()
print a_doubled
print a_gpu
```

Elementwise functions (e.g. `sqrt()`) are available in the module `pycuda.cumath`. Reductions can also be performed, e.g. `pycuda.gpuarray.sum(a)`.

Also the `pycuda.elementwise.ElementwiseKernel` class allows one to put in the relevant snippets of C code. Similarly there one can create a

custom *Map Reduce* function. See the documentation for details. Also note that when timing any kernel functions, be sure that the compilation has happened first.

## Random numbers

For performance, one should use the GPU optimized random number generators. The XORWOW generator should be good enough:

```
import pycuda.curandom
rg = pycuda.curandom.XORWOWRandomNumberGenerator()
arrayrand = rg.gen_uniform(100, numpy.float32)
```

`arrayrand` is now a GPUArray of random numbers.