Using GPUs on Hyak

Connect to hyak:

`ssh hyak.uw.edu`

Start an interactive session on a GPU node:

`qsub -I -q gpu -W group_list=hyak-hpc`

It can take a minute or so to get the prompt on the GPU node. Note that this gives you a session that lasts one hour by default.

(These are 12 core nodes (2 x 6 core Nehalem) with Tesla M2075 GPUs (Fermi, equivalent to GeForce GTX 480/590)

Set up your environment

Set up your environment to access the NVIDA Cuda compiler (`nvcc`):

`module load cuda_7`

For PyCUDA, set up your path to point at `/home/trq/python/bin`, and your `PYTHONPATH` to point at `/home/trq/python`. Edit your `.bashrc` to end with these lines:

```
export PATH="/usr/lusers/trq/python/bin:$PATH"
export PYTHONPATH="/usr/lusers/trq/python/lib/python:$PYTHONPATH"
```